

Assignment - TensorFlow and Keras Build various MLP architectures for MNIST dataset

In this assignment we have to try three different MLP architecture on MNIST data set. We have input as 784 and output as 10.

The task is to try 3 different MLP architecture, with Batch Normalization, Dropout as Regularizer, ReLU as Activation function and Adam as Optimizer.

We will try following architectures

Architecture 1 : MLP+ Adam + BN + ReLu + Dropout (0.5) + 2 Hidden Layer

Architecture 2 : MLP+ Adam + BN + ReLu + Dropout (0.3) + 3 Hidden Layer

Architecture 3 : MLP+ Adam + BN + ReLu + Dropout (0.2) + 5 Hidden Layer

Output: The expected output for this assignment is the plot between Train/test loss vs Epoch and Accuracy for all the models we train .

```
In [1]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
import warnings
warnings.filterwarnings('ignore')
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
```

Using TensorFlow backend.

```
In [2]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
```

```
plt.legend()
plt.grid()
fig.canvas.draw()
```

```
In [3]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2])
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
In [5]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [6]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

```
In [7]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms Lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255
```

```
X_train = X_train/255
X_test = X_test/255
```

```
In [8]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])
```

```
# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs
```

```
Y_train = np_utils.to_categorical(y_train, 10)
```

```
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

```
In [9]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####
```

```
# https://keras.io/activations/

# Activations can either be used through an Activation Layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
In [10]: output_dim = 10
         input_dim = X_train.shape[1]

         batch_size = 128
         nb_epoch = 20
```

Architecture 1 : MLP + Adam + BN + ReLu + Dropout (0.5) with 2 Hidden Layer (512 - 256)

```
In [11]: MLP_model_2Layers = Sequential()
         MLP_model_2Layers.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.0)
         MLP_model_2Layers.add(BatchNormalization())
         MLP_model_2Layers.add(Dropout(0.5))

         MLP_model_2Layers.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
         MLP_model_2Layers.add(BatchNormalization())
         MLP_model_2Layers.add(Dropout(0.5))

         MLP_model_2Layers.add(Dense(output_dim, activation='softmax'))
         MLP_model_2Layers.summary()
         MLP_model_2Layers.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
         history_2layers_b_d = MLP_model_2Layers.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=()
```

WARNING:tensorflow:From C:\Anaconda\lib\site-packages\tensorflow\python\ops\resource_variable_ops.py:435: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---|--------------|---------|
| dense_1 (Dense) | (None, 512) | 401920 |
| batch_normalization_1 (Batch Normalization) | (None, 512) | 2048 |
| dropout_1 (Dropout) | (None, 512) | 0 |
| dense_2 (Dense) | (None, 256) | 131328 |
| batch_normalization_2 (Batch Normalization) | (None, 256) | 1024 |
| dropout_2 (Dropout) | (None, 256) | 0 |
| dense_3 (Dense) | (None, 10) | 2570 |

Total params: 538,890

Trainable params: 537,354

Non-trainable params: 1,536

WARNING:tensorflow:From C:\Anaconda\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 126us/step - loss: 0.4514 - accuracy: 0.8629 - val_loss: 0.1454 - val_accuracy: 0.9548

Epoch 2/20

60000/60000 [=====] - 7s 124us/step - loss: 0.2115 - accuracy: 0.9352 - val_loss: 0.1080 - val_accuracy: 0.9652

Epoch 3/20

60000/60000 [=====] - 7s 119us/step - loss: 0.1638 - accuracy: 0.9498 - val_loss: 0.0898 - val_accuracy: 0.9713

Epoch 4/20

60000/60000 [=====] - 7s 118us/step - loss: 0.1388 - accuracy: 0.9567 - val_loss: 0.0823 - val_accuracy: 0.9740

Epoch 5/20

60000/60000 [=====] - 7s 119us/step - loss: 0.1222 - accuracy: 0.9626 - val_loss: 0.0767 - val_accuracy: 0.9758

Epoch 6/20

```

60000/60000 [=====] - 7s 121us/step - loss: 0.1068 - accuracy: 0.9665 - val_loss: 0.0737 - val_accuracy:
0.9767
Epoch 7/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1006 - accuracy: 0.9680 - val_loss: 0.0659 - val_accuracy:
0.9784
Epoch 8/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0936 - accuracy: 0.9702 - val_loss: 0.0639 - val_accuracy:
0.9805
Epoch 9/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0835 - accuracy: 0.9736 - val_loss: 0.0657 - val_accuracy:
0.9799
Epoch 10/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0800 - accuracy: 0.9736 - val_loss: 0.0588 - val_accuracy:
0.9812
Epoch 11/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0759 - accuracy: 0.9757 - val_loss: 0.0636 - val_accuracy:
0.9821
Epoch 12/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0702 - accuracy: 0.9775 - val_loss: 0.0558 - val_accuracy:
0.9825
Epoch 13/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0689 - accuracy: 0.9784 - val_loss: 0.0555 - val_accuracy:
0.9826
Epoch 14/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0644 - accuracy: 0.9795 - val_loss: 0.0578 - val_accuracy:
0.9829
Epoch 15/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0629 - accuracy: 0.9792 - val_loss: 0.0524 - val_accuracy:
0.9844
Epoch 16/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0583 - accuracy: 0.9812 - val_loss: 0.0539 - val_accuracy:
0.9828
Epoch 17/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0590 - accuracy: 0.9807 - val_loss: 0.0556 - val_accuracy:
0.9824
Epoch 18/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0515 - accuracy: 0.9830 - val_loss: 0.0558 - val_accuracy:
0.9833cy:
Epoch 19/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0520 - accuracy: 0.9831 - val_loss: 0.0543 - val_accuracy:
0.9838
Epoch 20/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0512 - accuracy: 0.9833 - val_loss: 0.0554 - val_accuracy:
0.9822

```

```

In [12]: score = MLP_model_2Layers.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])

```

```
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history_2layers_b_d.history['val_loss']
ty = history_2layers_b_d.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0554369732551655

Test accuracy: 0.982200026512146



Architecture 2 : MLP+ Adam + BN + ReLu + Dropout (0.3) with 3 Hidden Layer (512- 256 -128)

```
In [13]: MLP_model_3Layers = Sequential()
MLP_model_3Layers.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))
MLP_model_3Layers.add(BatchNormalization())
MLP_model_3Layers.add(Dropout(0.3))

MLP_model_3Layers.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))
MLP_model_3Layers.add(BatchNormalization())
MLP_model_3Layers.add(Dropout(0.3))
```



```

MLP_model_3Layers.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None)))
MLP_model_3Layers.add(BatchNormalization())
MLP_model_3Layers.add(Dropout(0.3))

MLP_model_3Layers.add(Dense(output_dim, activation='softmax'))
MLP_model_3Layers.summary()
MLP_model_3Layers.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history_b_d = MLP_model_3Layers.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test,

```

Model: "sequential_2"

| Layer (type) | Output Shape | Param # |
|---|--------------|---------|
| ===== | ===== | ===== |
| dense_4 (Dense) | (None, 512) | 401920 |
| batch_normalization_3 (Batch Normalization) | (None, 512) | 2048 |
| dropout_3 (Dropout) | (None, 512) | 0 |
| dense_5 (Dense) | (None, 256) | 131328 |
| batch_normalization_4 (Batch Normalization) | (None, 256) | 1024 |
| dropout_4 (Dropout) | (None, 256) | 0 |
| dense_6 (Dense) | (None, 128) | 32896 |
| batch_normalization_5 (Batch Normalization) | (None, 128) | 512 |
| dropout_5 (Dropout) | (None, 128) | 0 |
| dense_7 (Dense) | (None, 10) | 1290 |
| ===== | ===== | ===== |

Total params: 571,018

Trainable params: 569,226

Non-trainable params: 1,792

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 145us/step - loss: 0.3151 - accuracy: 0.9037 - val_loss: 0.1213 - val_accuracy: 0.9620

Epoch 2/20

60000/60000 [=====] - 8s 134us/step - loss: 0.1488 - accuracy: 0.9545 - val_loss: 0.0886 - val_accuracy: 0.9715

Epoch 3/20

60000/60000 [=====] - 8s 136us/step - loss: 0.1134 - accuracy: 0.9653 - val_loss: 0.0776 - val_accuracy:

```
0.9762
Epoch 4/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0976 - accuracy: 0.9704 - val_loss: 0.0827 - val_accuracy:
0.9742
Epoch 5/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0857 - accuracy: 0.9736 - val_loss: 0.0675 - val_accuracy:
0.9783
Epoch 6/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0727 - accuracy: 0.9768 - val_loss: 0.0704 - val_accuracy:
0.9808
Epoch 7/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0691 - accuracy: 0.9782 - val_loss: 0.0666 - val_accuracy:
0.9789
Epoch 8/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0619 - accuracy: 0.9799 - val_loss: 0.0625 - val_accuracy:
0.9804
Epoch 9/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0568 - accuracy: 0.9820 - val_loss: 0.0569 - val_accuracy:
0.9840
Epoch 10/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0526 - accuracy: 0.9831 - val_loss: 0.0622 - val_accuracy:
0.9811
Epoch 11/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0491 - accuracy: 0.9842 - val_loss: 0.0596 - val_accuracy:
0.9824
Epoch 12/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0455 - accuracy: 0.9858 - val_loss: 0.0610 - val_accuracy:
0.9825
Epoch 13/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0437 - accuracy: 0.9861 - val_loss: 0.0548 - val_accuracy:
0.9839
Epoch 14/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0428 - accuracy: 0.9859 - val_loss: 0.0612 - val_accuracy:
0.9818
Epoch 15/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0384 - accuracy: 0.9877 - val_loss: 0.0600 - val_accuracy:
0.9830
Epoch 16/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0367 - accuracy: 0.9882 - val_loss: 0.0552 - val_accuracy:
0.9835
Epoch 17/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0367 - accuracy: 0.9882 - val_loss: 0.0578 - val_accuracy:
0.9830
Epoch 18/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0359 - accuracy: 0.9885 - val_loss: 0.0591 - val_accuracy:
0.9843
Epoch 19/20
```

```
60000/60000 [=====] - 8s 131us/step - loss: 0.0323 - accuracy: 0.9897 - val_loss: 0.0592 - val_accuracy: 0.9836
Epoch 20/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0313 - accuracy: 0.9899 - val_loss: 0.0601 - val_accuracy: 0.9835
```

```
In [14]: score = MLP_model_3Layers.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history_b_d.history['val_loss']
ty = history_b_d.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.06012513352438982
Test accuracy: 0.9835000038146973
```



Architecture 3 : MLP+ Adam + BN + ReLu + Dropout (0.2) with 5 Hidden Layer (512- 256 -128 – 64 - 32)

```
In [15]: MLP_model_5Layers = Sequential()  
MLP_model_5Layers.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))  
MLP_model_5Layers.add(BatchNormalization())  
MLP_model_5Layers.add(Dropout(0.2))  
  
MLP_model_5Layers.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))  
MLP_model_5Layers.add(BatchNormalization())  
MLP_model_5Layers.add(Dropout(0.2))
```

```

MLP_model_5Layers.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None)))
MLP_model_5Layers.add(BatchNormalization())
MLP_model_5Layers.add(Dropout(0.2))

MLP_model_5Layers.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.102, seed=None)))
MLP_model_5Layers.add(BatchNormalization())
MLP_model_5Layers.add(Dropout(0.2))

MLP_model_5Layers.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.144, seed=None)))
MLP_model_5Layers.add(BatchNormalization())
MLP_model_5Layers.add(Dropout(0.2))

MLP_model_5Layers.add(Dense(output_dim, activation='softmax'))
MLP_model_5Layers.summary()
MLP_model_5Layers.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history_5layers_b_d = MLP_model_5Layers.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(

```

Model: "sequential_3"

| Layer (type) | Output Shape | Param # |
|---|--------------|---------|
| ===== | | |
| dense_8 (Dense) | (None, 512) | 401920 |
| batch_normalization_6 (Batch Normalization) | (None, 512) | 2048 |
| dropout_6 (Dropout) | (None, 512) | 0 |
| dense_9 (Dense) | (None, 256) | 131328 |
| batch_normalization_7 (Batch Normalization) | (None, 256) | 1024 |
| dropout_7 (Dropout) | (None, 256) | 0 |
| dense_10 (Dense) | (None, 128) | 32896 |
| batch_normalization_8 (Batch Normalization) | (None, 128) | 512 |
| dropout_8 (Dropout) | (None, 128) | 0 |
| dense_11 (Dense) | (None, 64) | 8256 |
| batch_normalization_9 (Batch Normalization) | (None, 64) | 256 |
| dropout_9 (Dropout) | (None, 64) | 0 |

| | | |
|--|------------|------|
| dense_12 (Dense) | (None, 32) | 2080 |
| batch_normalization_10 (Batch Normalization) | (None, 32) | 128 |
| dropout_10 (Dropout) | (None, 32) | 0 |
| dense_13 (Dense) | (None, 10) | 330 |

=====
 Total params: 580,778
 Trainable params: 578,794
 Non-trainable params: 1,984

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 10s 166us/step - loss: 0.4479 - accuracy: 0.8729 - val_loss: 0.1385 - val_accuracy: 0.9582

Epoch 2/20

60000/60000 [=====] - 9s 150us/step - loss: 0.1802 - accuracy: 0.9499 - val_loss: 0.1045 - val_accuracy: 0.9696

Epoch 3/20

60000/60000 [=====] - 9s 151us/step - loss: 0.1341 - accuracy: 0.9627 - val_loss: 0.0975 - val_accuracy: 0.9713

Epoch 4/20

60000/60000 [=====] - 9s 151us/step - loss: 0.1159 - accuracy: 0.9672 - val_loss: 0.0801 - val_accuracy: 0.9766

Epoch 5/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0986 - accuracy: 0.9722 - val_loss: 0.0839 - val_accuracy: 0.9768

Epoch 6/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0881 - accuracy: 0.9752 - val_loss: 0.0788 - val_accuracy: 0.9772

Epoch 7/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0832 - accuracy: 0.9764 - val_loss: 0.0723 - val_accuracy: 0.9787

Epoch 8/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0757 - accuracy: 0.9788 - val_loss: 0.0857 - val_accuracy: 0.9770

Epoch 9/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0678 - accuracy: 0.9803 - val_loss: 0.0665 - val_accuracy: 0.9812

Epoch 10/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0599 - accuracy: 0.9830 - val_loss: 0.0753 - val_accuracy: 0.9785

Epoch 11/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0594 - accuracy: 0.9824 - val_loss: 0.0671 - val_accuracy: 0.9821

Epoch 12/20

```

60000/60000 [=====] - 9s 147us/step - loss: 0.0554 - accuracy: 0.9836 - val_loss: 0.0676 - val_accuracy:
0.9814
Epoch 13/20
60000/60000 [=====] - 9s 150us/step - loss: 0.0507 - accuracy: 0.9855 - val_loss: 0.0672 - val_accuracy:
0.9816
Epoch 14/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0537 - accuracy: 0.9847 - val_loss: 0.0739 - val_accuracy:
0.9809
Epoch 15/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0490 - accuracy: 0.9858 - val_loss: 0.0710 - val_accuracy:
0.9800
Epoch 16/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0412 - accuracy: 0.9880 - val_loss: 0.0630 - val_accuracy:
0.9831
Epoch 17/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0413 - accuracy: 0.9876 - val_loss: 0.0653 - val_accuracy:
0.9834
Epoch 18/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0387 - accuracy: 0.9892 - val_loss: 0.0680 - val_accuracy:
0.9827
Epoch 19/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0396 - accuracy: 0.9887 - val_loss: 0.0695 - val_accuracy:
0.9824
Epoch 20/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0393 - accuracy: 0.9889 - val_loss: 0.0670 - val_accuracy:
0.9829

```

```

In [16]: score = MLP_model_5Layers.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss

```

```
# acc : train accuracy  
# for each key in history.history we will have a list of length equal to number of epochs  
  
vy = history_5layers_b_d.history['val_loss']  
ty = history_5layers_b_d.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06700490548466333

Test accuracy: 0.9829000234603882



Conclusion


```
In [17]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Sr. No", "Model", "Hidden Layer Info", "Test Accuracy (%)"]
x.add_row(["Architecture 1 ", "MLP+ Adam + BN + ReLu + Dropout (0.5) + 2 Layer ", "512 - 256", 98.22])
x.add_row(["Architecture 2 ", "MLP+ Adam + BN + ReLu + Dropout (0.3) + 3 Layer ", "512 - 256 -128" , 98.35])
x.add_row(["Architecture 3 ", "MLP+ Adam + BN + ReLu + Dropout (0.2) + 5 Layer ", "512- 256 -128 - 64 - 32", 98.29])
print(x)
```

| Sr. No | Model | Hidden Layer Info | Test Accuracy (%) |
|----------------|---|-------------------------|-------------------|
| Architecture 1 | MLP+ Adam + BN + ReLu + Dropout (0.5) + 2 Layer | 512 - 256 | 98.22 |
| Architecture 2 | MLP+ Adam + BN + ReLu + Dropout (0.3) + 3 Layer | 512 - 256 -128 | 98.35 |
| Architecture 3 | MLP+ Adam + BN + ReLu + Dropout (0.2) + 5 Layer | 512- 256 -128 - 64 - 32 | 98.29 |